

Public Case Study Series: From Intake to Verified Outcome

A sanitized first-party case on publishing a governed AI whitepaper without exposing the private blueprint

BRAIZ Works LLC · WP-09 · Version 1.0.0 · September 24, 2026

Public-safe research. Principles and control outcomes only; proprietary implementation remains protected.

Executive summary

This case study shows how a public research artifact can move from outcome definition through planning, assurance, publication, public readback, and closeout while preserving proprietary implementation boundaries. It uses the published WP-02 research release as a first-party example and distinguishes observed facts from modeled or illustrative claims.

Case boundary

This case covers the public release of BRAIZ Works Research WP-02, *The Enterprise Execution Gap*. The public release is suitable as a case because its repository, release identity, and reader surface were deliberately separated from protected UBuildOS implementation details. [7]

The case does not disclose private prompts, internal controller logic, attack sets, private evidence schemas, credentials, or protected recovery procedures.

Outcome

The intended outcome was not merely “write a whitepaper.” It was to create a public-safe research artifact that could serve community education and enterprise social proof while preserving first-party intellectual property boundaries. The publication needed consistent Markdown, PDF, and HTML representations; citation metadata; manifests and checksums; a public repository; a release identity; and public readback.

Observed lifecycle

The project separated planning, producer qualification, independent review, owner acceptance, freeze, publication authorization, deployment, public verification, social disposition, and final closeout. This prevented a successful build from being treated as automatic publication authority.

What the canary exposed

The publication canary surfaced several practical failure modes before the broader portfolio was batched: hidden-file upload friction, accidental path duplication, platform-specific Pages enablement, release-asset assembly, and the need to preserve exact owner authority for consequential lifecycle transitions.

Those lessons are valuable because they occurred at the operational seams rather than in the paper prose itself.

Observed versus modeled

Statement	Evidence class
WP-02 was published to a public GitHub repository and reader surface	Documented/observed first-party state
Exact public release files were hash-verified during the governed process	Documented first-party evidence

Statement	Evidence class
The workflow reduces future portfolio rework	Expected/operational inference, not a measured universal result
The same workflow will always prevent publication defects	Not claimed

Community value without blueprint leakage

The published paper explains enterprise assurance principles—authority, evidence, lifecycle, recovery, change control, and independent review—without exposing the exact system that implements them. That separation is the central design lesson for public technical documentation: publish the problem, principles, interfaces, and evidence of seriousness; protect the reconstruction-sensitive recipe.

Why the portfolio moved to batching

Once the WP-02 canary proved the publication mechanics, running nine more papers through nine fully serial pipelines would duplicate shared work. The successor portfolio plan therefore treats WP-02 as a closed canary and batches shared research, artifact generation, cross-paper QA, corpus leakage analysis, and assurance while preserving per-paper failure isolation.

Reusable case-study template

A public case study should record: exact outcome, source boundary, lifecycle path, evidence classes, material failures and repairs, measured versus modeled quantities, privacy/IP constraints, limitations, and what was deliberately not disclosed. That structure creates credibility without converting the case into a reconstruction guide.

Limitations

This is a first-party case and therefore should not be treated as an independent market study. It documents a bounded internal-to-public production workflow and does not establish general performance, cost, or acquisition outcomes for other organizations.

References

[1] Stanford Institute for Human-Centered Artificial Intelligence. *The 2026 AI Index Report — Economy*. 2026. <https://hai.stanford.edu/ai-index/2026-ai-index-report/economy>

[2] McKinsey & Company. *AI at work but not at scale*. December 10, 2025. <https://www.mckinsey.com/featured-insights/charts/ai-at-work-but-not-at-scale>

[3] Tabassi, E. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. National Institute of Standards and Technology, 2023. <https://doi.org/10.6028/NIST.AI.100-1>

[4] Autio, C., Schwartz, R., Dunietz, J., Jain, S., Stanley, M., Tabassi, E., Hall, P., Roberts, K. *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST AI 600-1, 2024;

NIST page updated April 8, 2026. <https://doi.org/10.6028/NIST.AI.600-1>

[5] National Institute of Standards and Technology. *NIST AI RMF Playbook*.
<https://www.nist.gov/itl/ai-risk-management-framework/nist-ai-rmf-playbook>

[6] BRAIZ Works LLC. *UBuildOS Verified Completion™ Thesis v1.0.5*. September 15, 2026. DOI:
<https://doi.org/10.5281/zenodo.22775537> ; public repository:
<https://github.com/BRAIZ-Works/ubuildos-verified-completion-thesis>

[7] BRAIZ Works LLC. *The Enterprise Execution Gap*. BRAIZ Works Research WP-02 v1.0.0. September 23, 2026.

Source-currentness note

External references [1]-[5] are carried forward from the independently reviewed and publicly released WP-02 v1.0.0 canary. This batch does not claim a new live-web currentness check. Time-sensitive statistics are used only where explicitly cited and dated. First-party references [6]-[7] are BRAIZ Works publications.

Rights and disclosure boundary

© 2026 BRAIZ Works LLC. All rights reserved.

This public artifact intentionally excludes private prompts, proprietary routing/controller logic, internal thresholds, private attack sets, protected evidence schemas, credentials, customer data, and reconstruction-sensitive implementation details.