

What Enterprises Should Demand From Agentic AI Systems

A buyer-facing control checklist for authority, evidence, recovery, lifecycle truth, and independent assurance

BRAIZ Works LLC · WP-10 · Version 1.0.0 · September 24, 2026

Public-safe research. Principles and control outcomes only; proprietary implementation remains protected.

Executive summary

Enterprise buyers should evaluate agentic AI systems on more than model capability and task automation. This paper presents an implementation-agnostic checklist of outcome-level controls that procurement, risk, engineering, security, and operations teams can require without demanding a vendor’s proprietary internal blueprint.

Capability is not the whole product

Model quality matters, but enterprise deployment adds requirements that benchmarks alone do not answer. Buyers need to know who can authorize consequential actions, how evidence is preserved, what happens when a workflow fails halfway through, how versions are controlled, and whether completion claims survive independent challenge.

Ten control families

Control family	Buyer question
Outcome contract	Can the system state exactly what “done” means?
Authority	Are reads, drafts, writes, sends, spend, and publication separately controlled?
Source/currentness	Can the system distinguish current authority from stale context?
Evidence	Are material claims attributable and recomputable?
Lifecycle	Does the system avoid calling candidates “complete”?
Change control	Does a material change invalidate affected proof?
Recovery	Are restores tested, not merely documented?
Independent review	Can a separate reviewer challenge exact subjects without repair?
Security/privacy	Are protected data and permissions fail-closed?
Observability/audit	Can material actions and decisions be reconstructed?

Demand exact authority boundaries

A buyer should be able to distinguish what the system may read, draft, modify, send, purchase, publish, deploy, or delete. Broad statements such as “the agent has access” are not enough for consequential workflows. Authority should be scoped to operation, subject, channel, and relevant lifecycle gate.

Demand evidence, not screenshots of success

Ask how the system binds evidence to exact versions, how stale evidence is invalidated, whether manifests and hashes are independently recomputed, and whether the reviewer can reproduce admission-critical facts without relying on private producer context.

Demand partial-failure semantics

Agentic systems interact with APIs and user interfaces that can fail after side effects occur. Buyers should ask what happens when an operation times out after a write, whether the system reads back state before retry, and how it prevents duplicate or destructive actions.

Demand lifecycle honesty

Procurement and risk teams should reject systems that blur “generated,” “tested,” “reviewed,” “approved,” “deployed,” and “observed.” Lifecycle truth is a control because downstream labels can otherwise hide uncompleted upstream obligations.

Demand recovery proof

Ask for evidence that backups are intact, restores have been tested, and the restored state is verified. Recovery should be part of the release posture for workflows that can materially affect enterprise data or operations.

Demand change-aware assurance

When prompts, tools, dependencies, models, policies, or code change, buyers should ask what proof is invalidated and what is rerun. A system that either reuses everything or retests everything has not demonstrated mature impact analysis.

Demand an independent challenge path

The producer should do deep self-testing, but final consequential assurance should have a separate review boundary. The buyer does not need the vendor’s private attack sets; it needs confidence that the reviewer is not merely restating the builder’s conclusion.

A procurement evidence request

For a consequential use case, request a bounded evidence package containing the exact subject identity, requirement/acceptance mapping, test results, recovery proof, change history, known limitations, lifecycle state, and independent-review disposition. The package should be portable enough that critical claims do not depend on inaccessible chat history.

Limitations

This checklist is not a compliance standard, certification, or legal procurement template. Applicability varies by risk, industry, jurisdiction, system architecture, and use case. Organizations should integrate these controls

with their own security, privacy, legal, regulatory, and operational requirements.

References

- [1] Stanford Institute for Human-Centered Artificial Intelligence. *The 2026 AI Index Report — Economy*. 2026. <https://hai.stanford.edu/ai-index/2026-ai-index-report/economy>
- [2] McKinsey & Company. *AI at work but not at scale*. December 10, 2025. <https://www.mckinsey.com/featured-insights/charts/ai-at-work-but-not-at-scale>
- [3] Tabassi, E. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. National Institute of Standards and Technology, 2023. <https://doi.org/10.6028/NIST.AI.100-1>
- [4] Autio, C., Schwartz, R., Dunietz, J., Jain, S., Stanley, M., Tabassi, E., Hall, P., Roberts, K. *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST AI 600-1, 2024; NIST page updated April 8, 2026. <https://doi.org/10.6028/NIST.AI.600-1>
- [5] National Institute of Standards and Technology. *NIST AI RMF Playbook*. <https://www.nist.gov/itl/ai-risk-management-framework/nist-ai-rmf-playbook>
- [6] BRAIZ Works LLC. *UBuildOS Verified Completion™ Thesis v1.0.5*. September 15, 2026. DOI: <https://doi.org/10.5281/zenodo.22775537> ; public repository: <https://github.com/BRAIZ-Works/ubuildos-verified-completion-thesis>
- [7] BRAIZ Works LLC. *The Enterprise Execution Gap*. BRAIZ Works Research WP-02 v1.0.0. September 23, 2026.

Source-currentness note

External references [1]-[5] are carried forward from the independently reviewed and publicly released WP-02 v1.0.0 canary. This batch does not claim a new live-web currentness check. Time-sensitive statistics are used only where explicitly cited and dated. First-party references [6]-[7] are BRAIZ Works publications.

Rights and disclosure boundary

© 2026 BRAIZ Works LLC. All rights reserved.

This public artifact intentionally excludes private prompts, proprietary routing/controller logic, internal thresholds, private attack sets, protected evidence schemas, credentials, customer data, and reconstruction-sensitive implementation details.